

Performance improvements of commonly used optimizers in machine learning using **CUDA** and **OpenCL**

Lennart Hahner

Table of Contents


- I. Motivation
- II. Related work & expectations
- III. Conceptual differences of OpenCL and CUDA
- IV. The implementation
- V. Results



Motivation


The role of optimization in machine learning

Optimizers are the engines of machine learning models.




The role of optimization in machine learning

Train learnable weights towards a minimum of a loss function.




The role of optimization in machine learning

Most inputs i.e. images or text are represented as matrices...




The role of hardware in machine learning

Thus we often deal with **matrix** or tensor operations during optimizations...




The role of hardware in machine learning

Thus we often deal with matrix or **tensor** operations during optimizations...




The role of hardware in machine learning

Matrix and Tensor operations are computationally expensive...



The role of hardware in machine learning

...exploiting hardware for parallelization





The role of hardware in machine learning

...exploiting hardware for parallelization

Specification	Intel Core i9 14900KF [3]	Nvidia GeForce RTX 5080 [4]
Cores	24	16384+
Max. Boost Clock (Ghz)	6	2.62
Architecture	x86	Blackwell
Core Paradigma	MIMD	SPMD

The role of hardware in machine learning

...how to exploit GPUs?






Expectations

Results of „From CUDA to OpenCL“ [2]


OpenCL enables code portability, but not performance



Results of „A comprehensive Performance comparision of CUDA and OpenCL“ [3]

CUDA is at most ~30% faster than OpenCL for most applications


$$PR = \frac{Perf(OpenCL)}{Perf(CUDA)}$$



Workloads

About more results

Most of the comparison results are more than 10 years old...



So what to expect from the comparison?


...CUDA might perform better, due to the longer compile/build time of OpenCL, but it's worth revisiting...




Conceptual comparison of CUDA and OpenCL



Abstraction differences




[8] Figure 1.7 in Aftab M. et al. "OpenCL Programming Guide"




[9] Figure 3.13 in Kirk D. et al. "Programming Massively Parallel Processors"



Architecture differences



[8] Figure 1.7 in Aftab M. et al. "OpenCL Programming Guide"




[9] Figure 3.13 in Kirk D. et al. "Programming Massively Parallel Processors"

Architecture differences

Unlike CUDA, OpenCL requires
an environmental setup...



1. Get `cl_device_id`
2. Create `cl_context`
3. Create `cl_command_queue`
4. Create `cl_program`
5. Define `cl_kernel`
6. Allocate `cl_mem`



Architecture differences

NVCC enables easy integration of
kernel functions...



1. Allocate device memory

`cudaMalloc(...)`


2. Copy to device `cudaMemcpy(...)`

3. Launch kernel

`kernel«blocks, threads»(...)`

4. Copy from device `cudaMemcpy(...)`

5. Free `cudaFree(...)`



Kernels in different languages



```
__global__ void adam_update_kernel_impl(  
    float *__restrict__ param,  
    const float *__restrict__ grad,  
    float *__restrict__ m,  
    float *__restrict__ v,  
    size_t n,  
    float lr, float beta1,  
    float beta2,  
    float bc1, float bc2, float eps)  
{  
    ...  
}
```




```
__kernel void adam_update(  
    __global float* param,  
    __global const float* grad,  
    __global float* m,  
    __global float* v,  
    const float lr,  
    const float beta1,  
    const float beta2,  
    const float eps,  
    const float bc1,  
    const float bc2,  
    const int n  
) {  
    ...  
}
```




The implementation

Architecture overview



Architecture overview



Benchmark layer

Definition and initialization of workloads...

Benchmark Module

Define the workload and run with optimizer



Workload Modules

<<Interface>>


workload

forward(...) and compute_loss(...)

GEMM


...

Architecture overview




Optimization layer

Optimizer implementations on different platforms




Architecture overview



Util layer


Implementation of cross-Cutting concerns



Results

Benchmark environment

```
__global__ void adam_update_kernel_impl(  
    float *__restrict__ param,  
    const float *__restrict__ grad,  
    float *__restrict__ m,  
    float *__restrict__ v,  
    size_t n,  
    float lr, float beta1,  
    float beta2,  
    float bc1, float bc2, float eps)  
{  
    ...  
}
```



Measured aspects in consideration

Metric	Measurement
Computation-Time	Measure start and end time difference of kernel execution.
Data-Transfer-time	Measure start and end time difference of host to device transfer.
Quality-of-solution	Compare the 10th batch run loss with expectation and normalize
Build-time	Measure time to build kernel


Benchmark environment

Running on...

A NVIDIA GeForce GTX 970
with an Intel i3-9100F and 16 GB RAM

Visuals of currently obtained measurements

Averaged computation time and data transfer time on batch-index....



Conclusion so far

OpenCL and CUDA perform on average pretty similar...

...but some Remarks:

- Currently only tested on GTX 970 → **old platform**
- The number of results are **sparse**
- Some metrics are missing, e.g. **build-time**

Outlook

Some TODOs until end of march...

- ✗ Platform independence

- ✗ AdamW

- ✗ More measurements

Sources

- [1] Rishi, Rahul. (2011). Improving the Character Recognition Efficiency of Feed Forward BP Neural Network. International Journal of Computer Science & Information Technology. 3. 10.5121/ijcsit.2011.3107.
- [2] Chi, Yujie & Schubert, Keith & Badal, Andreu & Roncali, Emilie. (2025). Review of GPU-based Monte Carlo simulation platforms for transmission and emission tomography in medicine. Physics in Medicine & Biology. 70. 10.1088/1361-6560/adfda7.
- [3] Intel, "Intel Core i9 Prozessor 14900KF", accessed on 05 Feb 2026 from "<https://www.intel.de/content/www/de/de/products/sku/236787/intel-core-i9-processor-14900kf-36m-cache-up-to-6-00-ghz/specifications.html>"
- [4] Nvidia, "Compare GeForce Graphics Cards" accessd on 05 Feb 2026 from <https://www.nvidia.com/en-us/geforce/graphics-cards/compare/#50-series>
- [5] „From CUDA to OpenCL: Towards a Performance-portable Solution for Multi-platform GPU Programming“ Peng D., Luszczek P., Tomov S., Peterson G., Dongarra J. Parallel Computing, Volume 38, Issue 8, Pages 391-407, 2012
- [6] J. Fang, A. L. Varbanescu and H. Sips, "A Comprehensive Performance Comparison of CUDA and OpenCL," 2011 International Conference on Parallel Processing, Taipei, Taiwan, 2012, pp. 216-225, doi: 10.1109/ICPP.2011.45.
- [7] Karimi, Kamran. "A Performance Comparison of CUDA and OpenCL." Arxiv Preprint ArXiv:1005.2581, arxiv.org, 2010.
- [8] Holm, Havard Heitlo & Brodtkorb, André & Sætra, Martin. (2020). Performance and Energy Efficiency of CUDA and OpenCL for GPU Computing Using Python. 10.3233/APC200089.
- [9] Puya Memarzia, Farshad Khunjush „An In-depth Study on the Performance Impact of CUDA, OpenCL, and PTX Code“ Journal of Information and Computing Science Vol. 10, No. 2, 2015, pp.124-136
- [10] Aftab M. et al. "OpenCL Programming Guide" 2012 Pearson Education, Inc. ISBN-13: 978-0-321-74964-2
- [11] Kirk D. et al. "Programming Massively Parallel Processors" 2010 David B. Kirk/NVIDIA Corporation and Wen-mei Hwu. Published by Elsevier Inc. ISBN: 978-0-12-381472-2